

KMeans clustering

Medical Dataset

Eric Yarger

```
In [1]: # Import Libraries
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
import missingno as msno
from scipy import stats
```

```
In [2]: # Windows 10, Anaconda, JupyterLab, JupyterNotebook
# Jupyter environment version
!jupyter --version

jupyter core      : 4.6.3
jupyter-notebook  : 6.0.3
qtconsole         : 4.7.2
ipython           : 7.13.0
ipykernel         : 5.1.4
jupyter client    : 6.1.2
jupyter lab       : 1.2.6
nbconvert         : 5.6.1
ipywidgets        : 7.5.1
nbformat          : 5.0.4
traitlets         : 4.3.3
```

```
In [3]: # Python Environment version
import platform
print(platform.python_version())

3.7.7
```

```
In [4]: df = pd.read_csv('C:/Users/ericy/Desktop/medical_clean.csv')
```

Dataset, Initial Investigation


```
In [5]: # Dataset, first glance
```

```
In [6]: df.isnull().sum()
```

```
Out[6]: CaseOrder      0
Customer_id      0
Interaction      0
UID              0
City             0
State            0
County           0
Zip              0
Lat              0
Lng              0
Population        0
Area             0
TimeZone         0
Job              0
Children         0
Age              0
Income           0
Marital          0
Gender           0
ReAdmis          0
VitD_levels      0
Doc_visits       0
Full_meals_eaten 0
vitD_supp        0
Soft_drink       0
Initial_admin    0
HighBlood        0
Stroke           0
Complication_risk 0
Overweight       0
Arthritis        0
Diabetes         0
Hyperlipidemia   0
BackPain         0
Anxiety          0
Allergic_rhinitis 0
Reflux_esophagitis 0
Asthma           0
Services         0
Initial_days     0
TotalCharge      0
Additional_charges 0
Item1            0
Item2            0
Item3            0
Item4            0
Item5            0
Item6            0
Item7            0
Item8            0
dtype: int64
```

```
In [7]: msno.matrix(df)
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1e7e60776c8>
```



```
In [8]: # Duplicate Check
df.duplicated().sum()
```

```
Out[8]: 0
```

```
In [9]: # Removing unnecessary columns
# Selecting ReAdmis (Target)
# Selecting continuous dependent features
dfs = df[['Income', 'Additional_charges']]
dfs
```


```
Out[9]:
```

	Income	Additional_charges
0	86575.93	17939.403420
1	46805.99	17612.998120
2	14370.14	17505.192460
3	39741.49	12993.437350
4	1209.56	3716.525786
...
9995	45967.61	8927.642000
9996	14983.02	28507.150000
9997	65917.81	15281.210000
9998	29702.32	7781.678000
9999	62682.63	11643.190000

10000 rows × 2 columns


```
In [10]: # Outlier Removal
# Boxplots, outlier detection
sns.boxplot(dfs['Income'])
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1e7e6532d08>
```




```
In [11]: sns.boxplot(dfs['Additional_charges'])
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1e7e65906c8>
```



```
In [12]: dfs.hist(figsize=(10,10))
```

```
Out[12]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001E7E7464788>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001E7E667DF08>]],
      dtype=object)
```



In [13]: *# Outlier removal method via Z-score, Code reference (Bushmanov, 2019)*


```
num_data = dfs.select_dtypes(include=['number'])
cat_data = dfs.select_dtypes(exclude=['number'])
```

In [14]: `idx = np.all(stats.zscore(num_data) < 3, axis=1)`

In [15]: `dfs = pd.concat([num_data.loc[idx], cat_data.loc[idx]], axis=1)`


In [16]: *# Post Outlier Removal*
Boxplots
`sns.boxplot(dfs['Income'])`

Out[16]: `<matplotlib.axes._subplots.AxesSubplot at 0x1e7e6972ac8>`



In [17]: `sns.boxplot(dfs['Additional_charges'])`

Out[17]: `<matplotlib.axes._subplots.AxesSubplot at 0x1e7e679ecc8>`



In [18]: dfs

Out[18]:

	Income	Additional_charges
0	86575.93	17939.403420
1	46805.99	17612.998120
2	14370.14	17505.192460
3	39741.49	12993.437350
4	1209.56	3716.525786
...
9995	45967.61	8927.642000
9996	14983.02	28507.150000
9997	65917.81	15281.210000
9998	29702.32	7781.678000
9999	62682.63	11643.190000

9857 rows × 2 columns

Step P2, Missing Data

In [19]: *# Check for and Handle any missing data*
dfs.isnull().sum()

Out[19]:

Income	0
Additional_charges	0
dtype:	int64

In [20]: dfs.isna().sum()

Out[20]:

Income	0
Additional_charges	0
dtype:	int64

In [21]: dfs.isnull().any()

Out[21]:

Income	False
Additional_charges	False
dtype:	bool

In [22]: dfs

Out[22]:

	Income	Additional_charges
0	86575.93	17939.403420
1	46805.99	17612.998120
2	14370.14	17505.192460
3	39741.49	12993.437350
4	1209.56	3716.525786
...
9995	45967.61	8927.642000
9996	14983.02	28507.150000
9997	65917.81	15281.210000
9998	29702.32	7781.678000
9999	62682.63	11643.190000

9857 rows × 2 columns

Step P3, Data Transformation

In [23]:

```
X = dfs
```

In [24]:

```
# Set Before Normalization
X
```

Out[24]:

	Income	Additional_charges
0	86575.93	17939.403420
1	46805.99	17612.998120
2	14370.14	17505.192460
3	39741.49	12993.437350
4	1209.56	3716.525786
...
9995	45967.61	8927.642000
9996	14983.02	28507.150000
9997	65917.81	15281.210000
9998	29702.32	7781.678000
9999	62682.63	11643.190000

9857 rows × 2 columns

In [25]:

```
# Data Normalization using Min/Max Scaling
# Code Reference (Data normalization with pandas, 2020)
X = (X - X.min()) / (X.max() - X.min())
```

In [26]:

```
# Set After Normalization
X
```

Out[26]:

	Income	Additional_charges
0	0.686851	0.539851
1	0.370773	0.527956
2	0.112984	0.524027
3	0.314627	0.359607
4	0.008389	0.021531
...
9995	0.364110	0.211438
9996	0.117855	0.924967
9997	0.522667	0.442979
9998	0.234839	0.169676
9999	0.496955	0.310400

9857 rows × 2 columns

In [27]:


```
# Check for Skew
# Code Reference (Python | Pandas dataframe.skew(), 2022)
```

```
X.skew(axis = 0, skipna = True)
```

```
Out[27]: Income          0.942369  
Additional_charges    0.831749  
dtype: float64
```

```
In [28]: X.hist(figsize=(10,10))
```

```
Out[28]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001E7E6838F48>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x000001E7E686F3C8>]],  
          dtype=object)
```



```
In [29]: X.mean()
```

```
Out[29]: Income          0.308010  
Additional_charges    0.357885  
dtype: float64
```

```
In [30]: X.median()
```

```
Out[30]: Income          0.263370  
Additional_charges    0.308104  
dtype: float64
```

```
In [31]: # Square-Root Transformation on dataset  
Xs = X.apply(np.sqrt)
```

```
In [32]: Xs.skew()
```

```
Out[32]: Income          0.201145  
Additional_charges    0.127993  
dtype: float64
```

```
In [33]: Xs.mean()
```



```
Out[33]: Income          0.524266  
Additional_charges    0.562884  
dtype: float64
```

```
In [34]: Xs.median()
```

```
Out[34]: Income          0.513196  
Additional_charges    0.555071  
dtype: float64
```

```
In [35]: Xs.hist(figsize=(10,10))
```

```
Out[35]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001E7E6946848>,  
                <matplotlib.axes._subplots.AxesSubplot object at 0x000001E7E69C04C8>]],  
          dtype=object)
```



In [36]: Xs

Out[36]:

	Income	Additional_charges
0	0.828765	0.734745
1	0.608912	0.726606
2	0.336131	0.723897
3	0.560916	0.599672
4	0.091589	0.146735
...
9995	0.603415	0.459824
9996	0.343301	0.961752
9997	0.722957	0.665567
9998	0.484602	0.411918
9999	0.704951	0.557135

9857 rows × 2 columns

In [37]: X = Xs

In [38]: `# Read out dataset`
`X.to_excel('C:/Users/ericy/Desktop/cleaned_data.xlsx')`

K-means clustering

In [39]: `# Prepared dataset cleaned_data`
`# Code reference (Sklearn.cluster.KMeans, n.d.)`
`import sklearn`

In [40]: `from sklearn.cluster import KMeans`

In [41]: X

Out[41]:


	Income	Additional_charges
0	0.828765	0.734745
1	0.608912	0.726606
2	0.336131	0.723897
3	0.560916	0.599672
4	0.091589	0.146735
...
9995	0.603415	0.459824
9996	0.343301	0.961752
9997	0.722957	0.665567
9998	0.484602	0.411918
9999	0.704951	0.557135

9857 rows × 2 columns

```
In [42]: # Initial KMeans. run for 1-11 clusters.
# Elbow Plot
#Code Reference (Machine learning - K-means, n.d.)
# Sum of square distance from cluster center
ssdistance = []

# Kmeans from 1 to 11 loop
for i in range(1,11):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(X)
    ssdistance.append(kmeans.inertia_)

plt.plot(range(1,11), ssdistance, marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('Distance to Cluster Center')
plt.show()
```



```
In [43]: from sklearn.metrics import silhouette_score
# Code Reference (Arvai, 2020)
```

```
In [44]: # KMeans clusters = 2
kmeans2 = KMeans(n_clusters = 2, init='k-means++', random_state=10)

# Fit the model
kmeans2.fit(X)

# Silhouette Score
kmeans2_silhouette = silhouette_score(X, kmeans2.labels_)

# Print Silhouette Score
print('Silhouette Score: %.3f' %kmeans2_silhouette)
```

Silhouette Score: 0.320

```
In [45]: # KMeans clusters = 3
kmeans3 = KMeans(n_clusters = 3, init='k-means++', random_state=10)

# Fit the model
kmeans3.fit(X)

# Silhouette Score
kmeans3_silhouette = silhouette_score(X, kmeans3.labels_)

# Print Silhouette Score
print('Silhouette Score: %.3f' %kmeans3_silhouette)
```

Silhouette Score: 0.352

```
In [46]: # KMeans clusters = 4
kmeans4 = KMeans(n_clusters = 4, init='k-means++', random_state=10)

# Fit the model
kmeans4.fit(X)

# Silhouette Score
kmeans4_silhouette = silhouette_score(X, kmeans4.labels_)

# Print Silhouette Score
print('Silhouette Score: %.3f' %kmeans4_silhouette)
```

Silhouette Score: 0.347

```
In [47]: # KMeans clusters = 5
kmeans5 = KMeans(n_clusters = 5, init='k-means++', random_state=10)

# Fit the model
kmeans5.fit(X)

# Silhouette Score
kmeans5_silhouette = silhouette_score(X, kmeans5.labels_)

# Print Silhouette Score
print('Silhouette Score: %.3f' %kmeans5_silhouette)
```

Silhouette Score: 0.332

```
In [48]: # KMeans clusters = 10
kmeans10 = KMeans(n_clusters = 10, init='k-means++', random_state=10)

# Fit the model
kmeans10.fit(X)

# Silhouette Score
kmeans10_silhouette = silhouette_score(X, kmeans10.labels_)

# Print Silhouette Score
print('Silhouette Score: %.3f' %kmeans10_silhouette)
```

Silhouette Score: 0.332

```
In [49]: import yellowbrick
from yellowbrick.cluster import SilhouetteVisualizer
#Code Reference (Yellowbrick :: Anaconda.org, n.d.)
```

```
In [50]: # Code Reference (Kumar, 2020)
# Silhouette score. Visualization using yellowbrick's Silhouette Visualizer

#plot sizing
fig, ax = plt.subplots(2, 2, figsize=(15,8))

#KMeans for 2,3,4,5 'k' clusters
for i in [2, 3, 4, 5]:


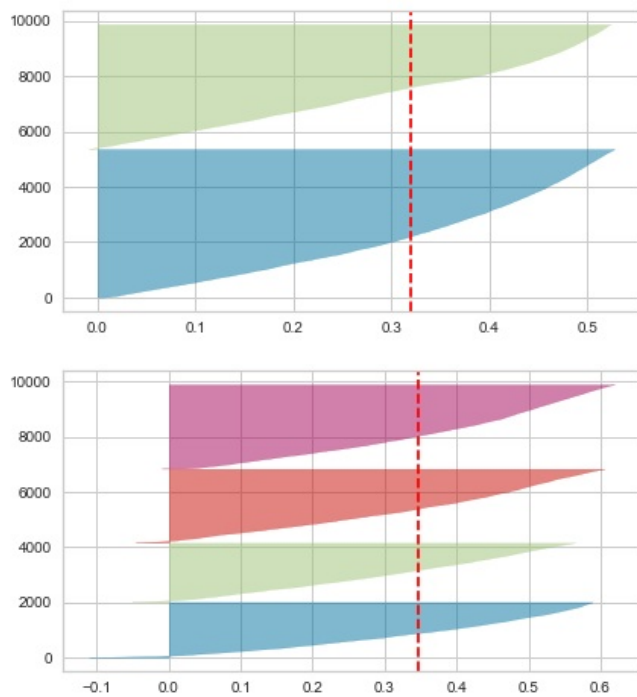
    km = KMeans(n_clusters=i, init='k-means++', n_init=10, max_iter=100, random_state=42)
    q, mod = divmod(i, 2)

    #Create SilhouetteVisualizer instance with KMeans instance

    visualization = SilhouetteVisualizer(km, colors='yellowbrick', ax=ax[q-1][mod])

    #Fit the visualizer

    visualization.fit(X)
```



In [51]: # Code Reference (How to plot KMeans?, n.d.)

```
centroids = kmeans3.cluster_centers_
labels = kmeans3.labels_


print(centroids)
print(labels)

[[0.51699125 0.81334543]
 [0.38289295 0.46267265]
 [0.70073493 0.46864602]]
[2 0 0 ... 2 1 2]
```

In [52]: colors = ["green", "red", "blue"]

```
plt.scatter(X.iloc[:,0], X.iloc[:,1], c=np.array(colors)[labels],
            s = 10, alpha=.1)

plt.scatter(centroids[:, 0], centroids[:, 1], marker = "x", s=150,
            linewidths = 5, zorder = 10, c=['green', 'red', 'blue'])
plt.show()
```



In [53]: X

Out[53]:

	Income	Additional_charges
--	--------	--------------------

0	0.828765	0.734745
1	0.608912	0.726606
2	0.336131	0.723897
3	0.560916	0.599672
4	0.091589	0.146735
...
9995	0.603415	0.459824
9996	0.343301	0.961752
9997	0.722957	0.665567
9998	0.484602	0.411918
9999	0.704951	0.557135

9857 rows × 2 columns

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js